

## PROJECT 1

### IMAGE THRESHOLDING AND BLOB TRACKING

#### OVERVIEW

For this project, you will build a system to:

- threshold a sequence of color or grayscale images to distinguish objects of interest from the background.
- apply morphological operators to the thresholded images in order to remove noise and imperfections.
- perform a connected components analysis to distinguish between separate objects, and to identify their locations in the image.
- track the connected components over time.

The last two items in the list above are collectively known as *blob tracking*.

#### TASKS

**Thresholding.** Your system will produce a thresholded binary image where the non-zero pixels correspond to objects of interest (foreground), and the zero pixels correspond to background. You will need to decide on some details of your thresholding approach:

*Averaging.* Instead of straightforwardly thresholding each frame, your system may need to perform some averaging in order to improve performance. We discussed two averaging strategies in class: spatial and temporal. In the spatial averaging approach, also known as *adaptive thresholding*, the system looks at the difference between a given pixel value and the average of the pixels around it. OpenCV implements adaptive thresholding with the **adaptiveThreshold** function. In the temporal averaging approach, the system looks at the difference between a given pixel value at some location and the average intensity of that location over time. You can take a temporal average by simply adding together multiple frames and dividing by the number of frames; however, be careful about overflow. In practice, this means you will want to convert frames from 8-bit integer format to a representation with more precision before summing them together. To perform a straightforward threshold on an image, you can use the OpenCV **threshold** function.

*RGB thresholding.* In class, we discussed a number of ways to threshold RGB-valued pixels. The simplest is to convert RGB to grayscale and threshold accordingly. Other methods include a planar decision boundary or distance from a reference RGB value. Although you may be tempted to write a loop over each pixel in the image, you should instead find some OpenCV or NumPy functionality to perform the equivalent operations more quickly.

*Intermediate image storage.* If you are averaging together multiple images, or subtracting images from each other, remember to first convert to a storage format that will allow these operations without overflow. Although your movie frames will be provided to you in `numpy.uint8` format, I would suggest using `numpy.float32` for intermediate images.

**Morphological operators.** Apply some morphological operators to the results of thresholding to eliminate noise and speckles. OpenCV provides the `erode`, `dilate`, and `morphologyEx` functions to implement erosion, dilation, and opening/closing, respectively. Your goal here is to produce the best possible image to send into the next stage of the processing pipeline.

**Connected components analysis.** The OpenCV function `findContours` retrieves the outlines of connected components of non-zero pixels. Applied to your thresholded image, this corresponds to outlines of the objects of interest in your scene. Some additional analysis of the contours yields information such as the area, centroid, and principal axes of each connected component (see the `regions.py` example from the course webpage for details).

**Tracking.** Your system should, at minimum, extract the position of each object's centroid in each frame. Better yet, it should also be able to track objects' trajectories over time by associating the connected components in the current frame with those of the previous frame (note that this is trivial in scenes containing a single object, but can become very tricky in scenes with many objects).

**Scenarios.** Your system should be targeted at some particular scenario. Here are some examples, ordered roughly in increasing difficulty:

- A single brightly colored object moving through the scene. Example: the video of the bright green cup that I showed in class.
- Multiple brightly colored objects moving through the scene. Example: video of someone juggling two or more differently colored balls.
- A single, arbitrarily colored object, tracked using temporal averaging. Example: the cat video I showed in class.
- Multiple objects tracked using temporal averaging. Example: the fruit flies video from class.

You should be careful to pick a scenario that is feasible given your programming ability and the time available. If you have questions about picking a particular scenario, or coming up with your own, don't hesitate to ask.

## EVALUATION CRITERIA

Your project will be evaluated by the following criteria (percentages are approximate):

- a. **source code and raw data (15%)** - Turn in the full source code for your project. I expect your code to be neatly indented and reasonably commented. Unless the data was provided by me, you should also turn at least two sets of raw data (movie files and any additional input necessary) that can be supplied to your program as input.
- b. **thresholding and morphological operators (25%)** - Your system should be able to output the thresholded binary images both before and after morphological operators have been applied. The output afterwards should be relatively free of noise and speckles, with the objects of interest well-distinguished from the background. Submit a set of several representative before/after image pairs from this intermediate data.
- c. **binary segmentation and tracking (25%)** - Produce plots of the positions of the objects of interest over time. At minimum, you should have a single  $(x, y)$  plot of disconnected points showing the positions of the centroids of the connected components in the thresholded images. Better yet, produce a plot of connected points showing the *trajectory* of each tracked object over time.
- d. **performance (5%)** - Your system should run in real-time or close to it. I expect it to be able to process  $640 \times 480$  video at least 10 frames per second (note that most video cameras generate output at about 30 fps). If you help downsampling videos (e.g. from a high-resolution camera), feel free ask how to do that on the course Piazza. Pre-processing for things like temporal averaging does not count against the real-time performance of your program.
- e. **“cool factor” (10%)** - Go a little above and beyond the tasks described in the section above. Possibilities include allowing parameters of the system to be modified interactively, visualizing some aspect of the data not described above, or doing something really creative or cool.
- f. **written report (20%)** - Turn in a written report in PDF format describing your overall approach and its effectiveness. The report should include the images and plots from **b** and **c**. In addition to addressing the above criteria, your report should also address some questions about generality. How broad a class of data can your system work on? What assumptions are encoded in your methods? Would your system work with objects of different colors? In changing brightness conditions?

Submit a zipfile containing your input videos, program source, and report to the course Moodle page. If your program creates video output (which should be relatively straightforward with `cv2.VideoWriter` – see `capture.py` in the example code), include it in your zipfile as well.